

Class 7: Machine Learning

Michelle Woo


Example of K-means clustering

First step is to make up some data with a known structure, so we know what the answer should be

```
# generating random data
tmp <- c(rnorm(30, mean=-3), rnorm(30, mean=3))
tmp
```

```
[1] -3.421605 -2.263818 -4.072744 -2.006747 -2.223410 -4.197326 -1.424704
[8] -3.580889 -3.633132 -1.978511 -2.212551 -2.790069 -2.393998 -2.481646
[15] -3.328053 -4.029252 -3.886351 -3.376245 -2.443714 -3.524116 -1.510642
[22] -1.282096 -3.554522 -2.486291 -2.876201 -3.712021 -3.907400 -4.123471
[29] -3.710735 -3.856055  5.542449  3.330610  3.774362  3.150696  2.874897
[36]  1.939550  4.079208  1.779662  1.822530  3.101260  4.157091  1.780063
[43]  2.096811  3.407007  3.503618  3.743822  2.596999  1.629745  4.049740
[50]  4.542132  2.739943  4.607806  3.346997  3.899235  2.854375  4.944517
[57]  3.147905  3.988761  2.485760  2.590351
```

```
# visualizing in 3D
x <- cbind(x=tmp, y=rev(tmp))
plot(x)
```



Now we have some structured data in `x`. Let's see if k-means is able to identify the two groups

```
k <- kmeans(x, centers=2, nst=20)
k
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	-3.009611	3.250263
2	3.250263	-3.009611

Clustering vector:

[illegible]

Within cluster sum of squares by cluster:

```
[1] 52.17816 52.17816
(between_SS / total_SS = 91.8 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

Let's explore and better understand k:

```
# how many elements are in each group?
k$size
```

[1] 30 30

k\$centers


	x	y
1	-3.009611	3.250263
2	3.250263	-3.009611

```
# able to use this to color the plot
k$cluster
```


[illegible]

Refining the plot:


```
# coloring the different groups
plot(x, col=k$cluster)
```



```
# adding in cluster centers, plot(x, col=k$cluster)
plot(k$centers)
```



```
# want to overlap the two above
plot(x, col=k$cluster)
points(k$centers, col = 'blue', pch = 8)
```



Example of Hierarchical Clustering

Let's use the same data as before, which we stored in `x`. We will use the `hclust()` function. `dist(x)` calculates the distance between all the points, this is input required for clustering

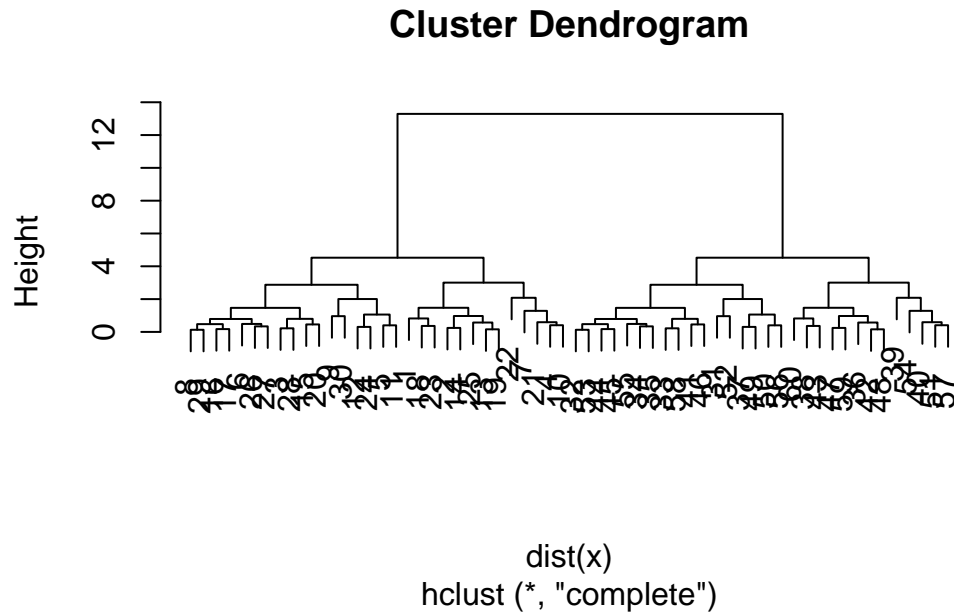
```
clustering <- hclust(dist(x))
clustering
```

Call:

```
hclust(d = dist(x))
```

```
Cluster method   : complete
Distance         : euclidean
Number of objects: 60
```


```
# results in tree, plot function gives something different
plot(clustering)
```



Lets add a horizontal line

```
plot(clustering)
abline(h=10, col='pink') # results in 6 classifications
```

Cluster Dendrogram



```
dist(x)
hclust (*, "complete")
```


To get our results (i.e. membership vector) we need to ‘cut’ the tree at the chosen height. The function for doing that is `cuttree()`

```
# able to get membership clustering
subgroups <- cutree(clustering, h=10)
subgroups
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

Plotting this:

```
plot(x, col= subgroups)
```



You can cut your tree with the number of clusters you want:

```
cutree(clustering, k=2)
```

[illegible]

Principal Component Analysis (PCA)

PCA of UK food

First, we need to read the data

```
url <- "https://tinyurl.com/UK-foods"

# making sure foods are the first column and for our rows
x <- read.csv(url, row.names=1)
head(x)
```


	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
nrow(x)
```

```
[1] 17
```

```
ncol(x)
```

```
[1] 4
```

```
# using dim() which provides both information
dim(x)
```

```
[1] 17 4
```

There are 17 rows and 5 columns in the data frame (after setting the correct row names)

Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

```
x <- read.csv(url, row.names=1)
```


```
rownames(x) <- x[,1]
x <- x[,-1]
head(x)
```

	Wales	Scotland	N.Ireland
105	103	103	66
245	227	242	267
685	803	750	586
147	160	122	93
193	235	184	209
156	175	147	139

I prefer the first method where it is just one line of code to correct the row names. It is more seamless and effortless. And if the second code block is run multiple times, it becomes an error as it removes each column with each run.

Now we can generate some basic visualizations. We need to make x as a matrix to be able to plot it


```
barplot(as.matrix(x))
```



```
rainbow(nrow(x))
```


```
[1] "#FF0000" "#FF5A00" "#FFB400" "#F0FF00" "#96FF00" "#3CFF00" "#00FF1E"
[8] "#00FF78" "#00FFD2" "#00D2FF" "#0078FF" "#001EFF" "#3C00FF" "#9600FF"
[15] "#F000FF" "#FF00B4" "#FF005A"
```

```
# combining - giving color to the plot  
barplot(as.matrix(x), col=rainbow(nrow(x)))
```



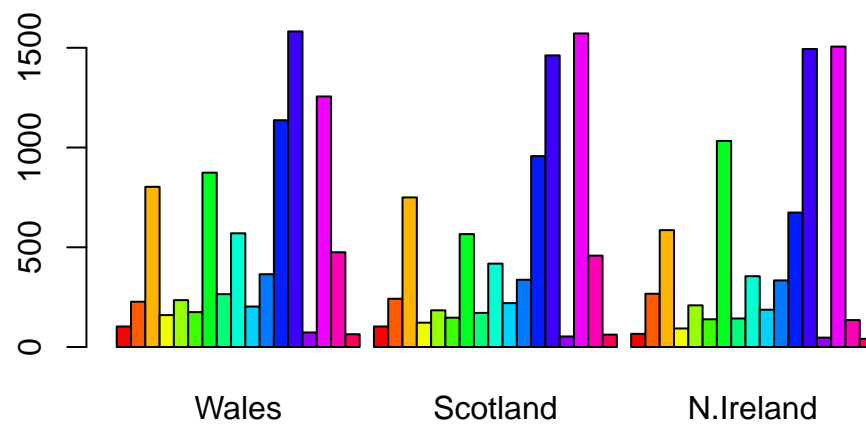
Lets refine our barplot

```
barplot(as.matrix(x), col=rainbow(nrow(x)), beside = T)
```




Q3. Changing what optional argument in the above `barplot()` function results in the following plot?

```
# beside to either T or F
barplot(as.matrix(x), col=rainbow(nrow(x)), beside = T)
```



```
barplot(as.matrix(x), col=rainbow(nrow(x)), beside = F)
```




Adding the argument `(beside = T)` to `barplot(as.matrix(x), col=rainbow(nrow(x)))` will result in the data to be stacked 'beside' each other. Removing that would result in the barplot with stacked data.

Q4. Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

If a point were to be shifted more towards one of the locations, that means that data point belongs to that area. Though the figure is difficult to discern and analyze.

Other visualizations that can be useful:

```
pairs(x, col=rainbow(nrow(x)), pch=18)
```



Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

N. Ireland consumes more fresh potatoes and soft drinks and less of the other foods hence its location at the bottom of the figure and the subsequent data points near it.

Lets apply PCA. For that, we need to use the command `prcomp()`. This function expects the transpose of our data

```
# t flips the rows and columns
# transpose_matrix <- t(x)
# pca <- prcomp(transpose_matrix)
pca <- prcomp(t(x))
```


```
summary(pca)
```

Importance of components:

	PC1	PC2	PC3
Standard deviation	379.8991	260.5533	1.515e-13
Proportion of Variance	0.6801	0.3199	0.000e+00
Cumulative Proportion	0.6801	1.0000	1.000e+00

Lets plot the PCA results:

```
plot(pca)
```



We need to access the results of the PCA analysis

```
attributes(pca)
```

```
$names
[1] "sdev"      "rotation" "center"    "scale"     "x"

$class
[1] "prcomp"
```

We can explore the `pca$x` dataframe:

(all 4 components, we can now place 2 in x axis and 2 in y axis)

```
pca$x
```


	PC1	PC2	PC3
Wales	-288.9534	226.36855	2.296774e-14
Scotland	-141.3603	-284.81172	4.517428e-13
N.Ireland	430.3137	58.44317	-1.407069e-13

Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.


Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.

Plotting:

```
plot(x=pca$x[,1], y=pca$x[,2])
```

```
# overlay country names and adding colors
plot(pca$x[,1], pca$x[,2] )
colors_countries <- c('orange', 'pink', 'blue', 'green')
text(x=pca$x[,1], y=pca$x[,2], colnames(x), col=colors_countries)
```



Plot of variances and loading scores

```
# calculating how much variation each PC accounts for
v <- round( pca$sdev^2/sum(pca$sdev^2) * 100 )
v
```


```
[1] 68 32 0
```

```
z <- summary(pca)
z$importance
```

	PC1	PC2	PC3
Standard deviation	379.89908	260.55330	1.515317e-13
Proportion of Variance	0.68009	0.31991	0.000000e+00
Cumulative Proportion	0.68009	1.00000	1.000000e+00


Plotting:

```
barplot(v, xlab="Principal Component", ylab="Percent Variation")
```



Making it into a biplot:


```
# Focusing on PC1 as it accounts for > 90% of variance  
par(mar=c(10, 3, 0.35, 0))  
barplot( pca$rotation[,1], las=2 )
```



Q9. Generate a similar 'loadings plot' for PC2. What two food groups feature prominently and what does PC2 mainly tell us about?

```
v <- round( pca$sdev^2/sum(pca$sdev^2) * 100 )

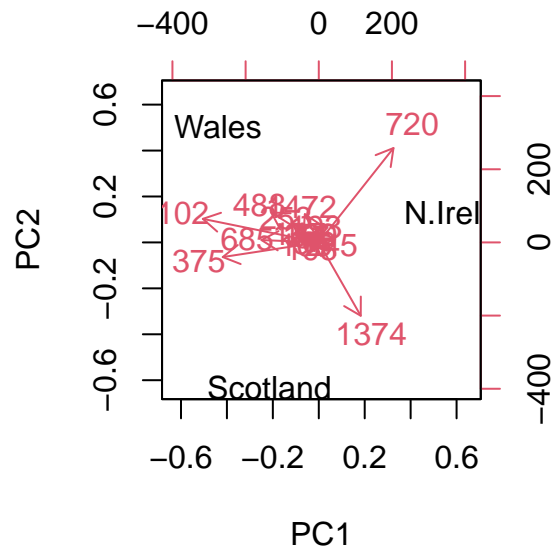
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,2], las=2 )
```



Fresh potatoes and soft drinks are the two main food groups featured prominently. PC2 would mainly tell us about the push of N. Ireland to the right positive side of the plot and the push of Scotland to the left negative side for soft drinks.

Biplots:

```
biplot(pca)
```



PCA of RNA-seq dataset

First step as always is to load the data:

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

	wt1	wt2	wt3	wt4	wt5	ko1	ko2	ko3	ko4	ko5
gene1	439	458	408	429	420	90	88	86	90	93
gene2	219	200	204	210	187	427	423	434	433	426
gene3	1006	989	1030	1017	973	252	237	238	226	210
gene4	783	792	829	856	760	849	856	835	885	894
gene5	181	249	204	244	225	277	305	272	270	279
gene6	460	502	491	491	493	612	594	577	618	638

Q10. How many genes and samples are in this data set?

```
dim(rna.data)
```

```
[1] 100 10
```

There are 100 genes, and 10 samples.

Now lets apply PCA:

```
pca_rna <- prcomp(t(rna.data))
summary(pca_rna)
```


Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6
Standard deviation	2214.2633	88.9209	84.33908	77.74094	69.66341	67.78516
Proportion of Variance	0.9917	0.0016	0.00144	0.00122	0.00098	0.00093
Cumulative Proportion	0.9917	0.9933	0.99471	0.99593	0.99691	0.99784

	PC7	PC8	PC9	PC10
Standard deviation	65.29428	59.90981	53.20803	3.142e-13
Proportion of Variance	0.00086	0.00073	0.00057	0.000e+00
Cumulative Proportion	0.99870	0.99943	1.00000	1.000e+00

Lets plot the principal components 1 and 2

```
plot(pca_rna$x[,1], pca_rna$x[,2], xlab='PC1', ylab='PC2')
```




```
# checking the names
colnames(rna.data)
```

```
[1] "wt1" "wt2" "wt3" "wt4" "wt5" "ko1" "ko2" "ko3" "ko4" "ko5"
```

```
# generating a vector that will color the sample
cols_samples <- c(rep('blue', 5), rep('red', 5))
cols_samples
```

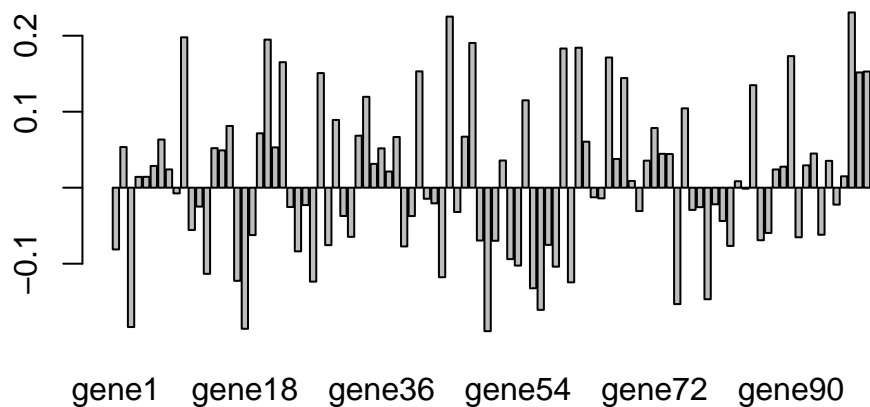
```
[1] "blue" "blue" "blue" "blue" "blue" "red" "red" "red" "red" "red"
```

```
# applying color to the plot
plot(pca_rna$x[,1], pca_rna$x[,2], xlab='PC1', ylab='PC2', col=cols_samples)
```



Identifying which gene is contributing the most

```
barplot(pca_rna$rotation[,1])
```

```
# identifying under- or overexpression
sort(pca_rna$rotation[,1])
```

gene50	gene18	gene3	gene57	gene75	gene79
-0.188796985	-0.185668500	-0.183374164	-0.160771014	-0.153164404	-0.146803635
gene56	gene61	gene27	gene17	gene44	gene13
-0.132330117	-0.124572881	-0.123615228	-0.122536548	-0.117808971	-0.113357525
gene59	gene54	gene53	gene25	gene1	gene39
-0.103935563	-0.102503320	-0.093979884	-0.083761992	-0.081247810	-0.077306742
gene82	gene29	gene58	gene51	gene49	gene86
-0.076658760	-0.075605635	-0.075274651	-0.069855142	-0.069530208	-0.069165267
gene91	gene32	gene19	gene94	gene87	gene11
-0.065288752	-0.064721235	-0.062411218	-0.061938300	-0.059547317	-0.055698801
gene81	gene40	gene31	gene46	gene70	gene77
-0.043780416	-0.037323670	-0.037219970	-0.031990529	-0.030784982	-0.029225446
gene78	gene24	gene12	gene26	gene96	gene80
-0.025639741	-0.025407507	-0.024870802	-0.022868107	-0.022293151	-0.021824860
gene43	gene42	gene65	gene64	gene9	gene84
-0.020617052	-0.014550791	-0.014052839	-0.012639567	-0.007495075	-0.001289937
gene83	gene69	gene4	gene5	gene97	gene37
0.008504287	0.008871890	0.014242602	0.014303808	0.014994546	0.021280555
gene88	gene8	gene89	gene6	gene92	gene35

0.024015925	0.024026657	0.027652967	0.028634131	0.029394259	0.031349942
gene95	gene71	gene52	gene67	gene74	gene73
0.035342407	0.035589259	0.035802086	0.037840851	0.044286948	0.044581700
gene93	gene15	gene36	gene14	gene22	gene2
0.044940861	0.049090676	0.051765605	0.052004194	0.053013523	0.053465569
gene63	gene7	gene38	gene47	gene33	gene20
0.060529157	0.063389255	0.066665407	0.067141911	0.068437703	0.071571203
gene72	gene16	gene30	gene76	gene55	gene34
0.078551648	0.081254592	0.089150461	0.104435777	0.114988217	0.119604059
gene85	gene68	gene28	gene99	gene100	gene41
0.134907896	0.144227333	0.150812015	0.151678253	0.152877246	0.153077075
gene23	gene66	gene90	gene60	gene62	gene48
0.165155192	0.171311307	0.173156806	0.183139926	0.184203008	0.190495289
gene21	gene10	gene45	gene98		
0.194884023	0.197905454	0.225149201	0.230633225		